

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import missingno as msno

import plotly.offline as py
py.init_notebook_mode(connected=True)

import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder

```

Dataset Describe

```

import pandas as pd

# Load the dataset to understand its structure and content
file_path = 'churn.csv'
data = pd.read_csv(file_path)

# Displaying the first few rows of the dataset
data.head()

```

	customer_id	credit_score	country	gender	age	tenure	balance	products_number	credit_card	active_member	estimated_salary	churn
0	15634602	619	France	Female	42	2	0.00	1	0	0	0	0
1	15647311	608	Spain	Female	41	1	83807.86	1	0	0	0	0
2	15619304	502	France	Female	42	8	159660.80	3	0	0	0	0
3	15701354	699	France	Female	39	1	0.00	2	0	0	0	0
4	15737888	850	Spain	Female	43	2	125510.82	1	0	0	0	0

```
data.describe().transpose()
```

	count	mean	std	min	25%	50%	75%	max
customer_id	10000.0	1.569094e+07	71936.186123	15565701.00	15628528.25	15690744.00	15737888.00	15737888.00
credit_score	10000.0	6.505288e+02	96.653299	350.00	584.00	652.00	699.00	850.00
age	10000.0	3.892180e+01	10.487806	18.00	32.00	37.00	43.00	43.00
tenure	10000.0	5.012800e+00	2.892174	0.00	3.00	5.00	8.00	8.00
balance	10000.0	7.648589e+04	62397.405202	0.00	0.00	97198.54	159660.80	159660.80
products_number	10000.0	1.530200e+00	0.581654	1.00	1.00	1.00	3.00	3.00
credit_card	10000.0	7.055000e-01	0.455840	0.00	0.00	1.00	1.00	1.00
active_member	10000.0	5.151000e-01	0.499797	0.00	0.00	1.00	1.00	1.00
estimated_salary	10000.0	1.000902e+05	57510.492818	11.58	51002.11	100193.90	157378.88	157378.88
churn	10000.0	2.037000e-01	0.402769	0.00	0.00	0.00	1.00	1.00

handling missing values

```
# Descriptive statistical analysis
descriptive_stats = data.describe()

# Checking for missing values
missing_values = data.isnull().sum()

descriptive_stats, missing_values

(
  customer_id  credit_score      age      tenure      balance \
count  1.000000e+04  10000.000000  10000.000000  10000.000000  10000.000000
mean    1.569094e+07    650.528800    38.921800    5.012800    76485.889288
std     7.193619e+04    96.653299    10.487806    2.892174    62397.405202
min     1.556570e+07    350.000000    18.000000    0.000000    0.000000
25%     1.562853e+07    584.000000    32.000000    3.000000    0.000000
50%     1.569074e+07    652.000000    37.000000    5.000000    97198.540000
75%     1.575323e+07    718.000000    44.000000    7.000000   127644.240000
max     1.581569e+07    850.000000    92.000000   10.000000   250898.090000

  products_number  credit_card  active_member  estimated_salary \
count    10000.000000  10000.000000  10000.000000  10000.000000
mean         1.530200    0.705500    0.515100    100090.239881
std          0.581654    0.455840    0.499797    57510.492818
min          1.000000    0.000000    0.000000    11.580000
25%          1.000000    0.000000    0.000000    51002.110000
50%          1.000000    1.000000    1.000000   100193.915000
75%          2.000000    1.000000    1.000000   149388.247500
max          4.000000    1.000000    1.000000   199992.480000

  churn
count  10000.000000
mean    0.203700
std     0.402769
min     0.000000
25%     0.000000
50%     0.000000
75%     0.000000
max     1.000000
customer_id  0
credit_score  0
country      0
gender       0
age          0
tenure       0
balance      0
products_number  0
credit_card   0
active_member  0
estimated_salary  0
churn         0
dtype: int64)
```

EDA

```
data.describe().T.sort_values(ascending = 0,by = "mean").style.background_gradient(cmap = "BuGn")\
.bar(subset = ["std"], color = "red").bar(subset = ["mean"], color = "blue")
```

	count	mean	std	min	
customer_id	10000.000000	15690940.569400	71936.186123	15565701.000000	15628528.2
estimated_salary	10000.000000	100090.239881	57510.492818	11.580000	51002.1
balance	10000.000000	76485.889288	62397.405202	0.000000	0.0
credit_score	10000.000000	650.528800	96.653299	350.000000	584.0
age	10000.000000	38.921800	10.487806	18.000000	32.0
tenure	10000.000000	5.012800	2.892174	0.000000	3.0
products_number	10000.000000	1.530200	0.581654	1.000000	1.0
credit_card	10000.000000	0.705500	0.455840	0.000000	0.0
active_member	10000.000000	0.515100	0.499797	0.000000	0.0
churn	10000.000000	0.203700	0.402769	0.000000	0.0

```
import matplotlib.pyplot as plt
import seaborn as sns

# Setting the aesthetic style of the plots
sns.set(style="whitegrid")

# Creating a figure with multiple subplots
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(15, 10))

# Plotting distributions of key variables
sns.histplot(data['credit_score'], kde=True, ax=axes[0, 0])
axes[0, 0].set_title('Distribution of Credit Score')

sns.histplot(data['age'], kde=True, ax=axes[0, 1])
axes[0, 1].set_title('Distribution of Age')

sns.histplot(data['balance'], kde=True, ax=axes[1, 0])
axes[1, 0].set_title('Distribution of Balance')

sns.histplot(data['estimated_salary'], kde=True, ax=axes[1, 1])
axes[1, 1].set_title('Distribution of Estimated Salary')

plt.tight_layout()
plt.show()

# Creating boxplots to identify outliers
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(15, 10))

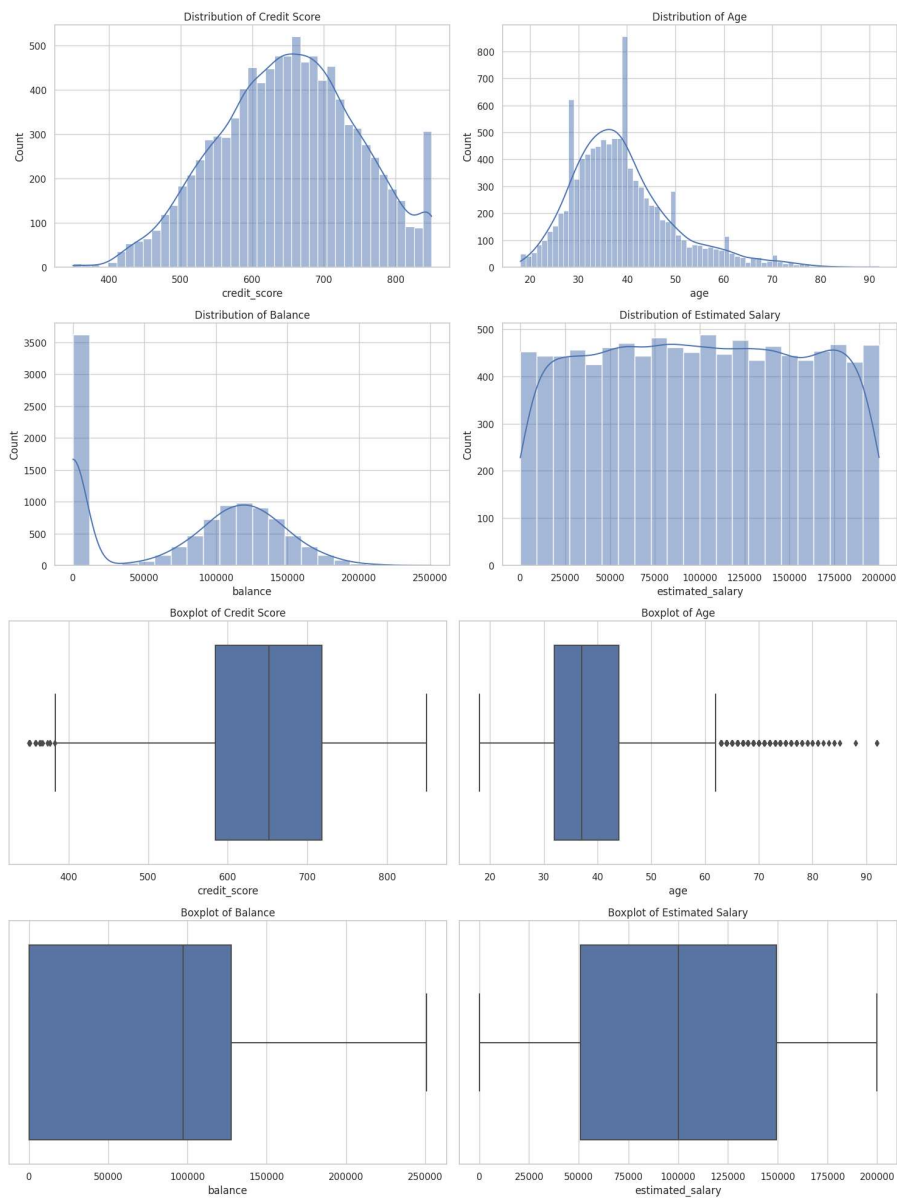
sns.boxplot(x='credit_score', data=data, ax=axes[0, 0])
axes[0, 0].set_title('Boxplot of Credit Score')

sns.boxplot(x='age', data=data, ax=axes[0, 1])
axes[0, 1].set_title('Boxplot of Age')

sns.boxplot(x='balance', data=data, ax=axes[1, 0])
axes[1, 0].set_title('Boxplot of Balance')

sns.boxplot(x='estimated_salary', data=data, ax=axes[1, 1])
axes[1, 1].set_title('Boxplot of Estimated Salary')

plt.tight_layout()
plt.show()
```



```
fig = px.histogram(data, x="age", y="balance", color="churn", marginal="box", hover_data=data.columns)
fig.show()
```

```
def plot_correlation_heatmap(data):
    numeric_data = data.select_dtypes(include=[np.number])
    plt.figure(figsize=(10, 8))
    sns.heatmap(numeric_data.corr(), annot=True, cmap='coolwarm')
    plt.title('Correlation Heatmap')
    plt.show()

def plot_categorical_churn(data):
    fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(15, 10))

    sns.countplot(x='country', hue='churn', data=data, ax=axes[0, 0])
    axes[0, 0].set_title('Churn by Country')

    sns.countplot(x='gender', hue='churn', data=data, ax=axes[0, 1])
    axes[0, 1].set_title('Churn by Gender')

    sns.countplot(x='products_number', hue='churn', data=data, ax=axes[1, 0])
    axes[1, 0].set_title('Churn by Number of Products')

    sns.countplot(x='credit_card', hue='churn', data=data, ax=axes[1, 1])
    axes[1, 1].set_title('Churn by Credit Card Possession')

    plt.tight_layout()
    plt.show()

def plot_numerical_scatter(data):
    fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(15, 5))

    sns.scatterplot(x='age', y='balance', hue='churn', data=data, ax=axes[0])
    axes[0].set_title('Age vs. Balance')







    sns.scatterplot(x='credit_score', y='estimated_salary', hue='churn', data=data, ax=axes[1])
    axes[1].set_title('Credit Score vs. Estimated Salary')

    plt.tight_layout()
    plt.show()
```

```
def plot_churn_distribution(data):
    churn_counts = [data[data['churn'] == 1].shape[0], data[data['churn'] == 0].shape[0]]
    churn_labels = ['Churned Customers', 'Retained Customers']
    fig, ax = plt.subplots(figsize=(10, 8))
    ax.pie(churn_counts, labels=churn_labels, shadow=True, autopct='%1.2f%%')
    plt.legend()
    plt.title("Comparison of Churned and Retained Customers", size=15)
    plt.show()

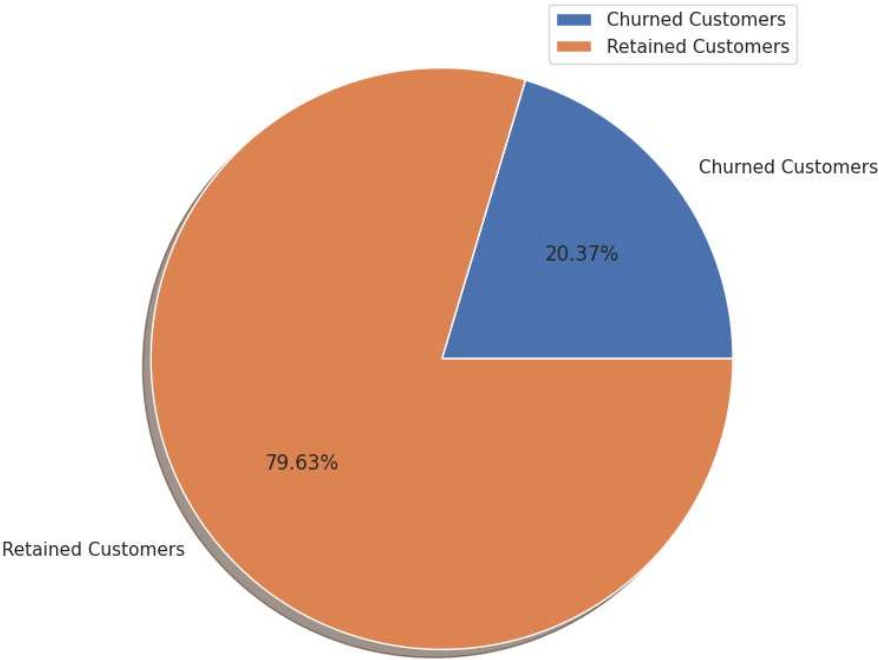
def display_grouped_statistics(data):
    numeric_columns = data.select_dtypes(include=[np.number]).columns
    feature_columns = numeric_columns.drop('churn') if 'churn' in numeric_columns else numeric_columns
    grouped_mean = data.groupby('churn')[feature_columns].mean().style.background_gradient(cmap="cool")
    grouped_median = data.groupby('churn')[feature_columns].median().style.background_gradient(cmap="cool")
    return grouped_mean, grouped_median

plot_correlation_heatmap(data)
plot_categorical_churn(data)
plot_numerical_scatter(data)
```



```
plot_churn_distribution(data)
mean_display, median_display = display_grouped_statistics(data)
mean_display
```

Comparison of Churned and Retained Customers



	customer_id	credit_score	age	tenure	balance	products_number
churn						
0	15691167.881703	651.853196	37.408389	5.033279	72745.296779	1.544267
1	15690051.964654	645.351497	44.837997	4.932744	91108.539337	1.475209

median_display

	customer_id	credit_score	age	tenure	balance	products_number
churn						
0	15691543.000000	653.000000	36.000000	5.000000	92072.680000	2.000000
1	15688963.000000	646.000000	45.000000	5.000000	109349.290000	1.000000


```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

# Defining the features and target variable
features = data.drop(['churn', 'customer_id'], axis=1) # Dropping 'customer_id' as it's not useful for prediction
target = data['churn']

# Identifying categorical and numeric columns
categorical_cols = features.select_dtypes(include=['object', 'category']).columns
numeric_cols = features.select_dtypes(include=['int64', 'float64']).columns

# Creating transformers for numeric and categorical columns
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Imputing missing values if any
    ('scaler', StandardScaler())]) # Scaling numeric features

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Imputing missing values if any
    ('onehot', OneHotEncoder(handle_unknown='ignore'))]) # One-hot encoding for categorical variables

# Combining transformers into a ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_cols),
        ('cat', categorical_transformer, categorical_cols)])

# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Applying the transformations to the training data
X_train = preprocessor.fit_transform(X_train)
X_test = preprocessor.transform(X_test)

X_train.shape, X_test.shape, y_train.shape, y_test.shape

((8000, 13), (2000, 13), (8000,), (2000,))

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Initializing the Logistic Regression model
logistic_model = LogisticRegression(random_state=42)

# Training the model
logistic_model.fit(X_train, y_train)

# Predicting on the test set
y_pred = logistic_model.predict(X_test)

# Calculating accuracy and other metrics
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(accuracy)

0.811

conf_matrix

array([[1543, 64],
       [ 314, 79]])

print(class_report)

              precision    recall  f1-score   support

0               0.83         0.96         0.89         1607
1               0.55         0.20         0.29          393

accuracy               0.81         2000

```

macro avg	0.69	0.58	0.59	2000
weighted avg	0.78	0.81	0.77	2000

```

from sklearn.metrics import roc_curve, auc

# Calculating the probabilities of the predictions
y_prob = logistic_model.predict_proba(X_test)[:, 1]


# Generating ROC curve values: fpr, tpr, thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob)

# Calculating AUC
roc_auc = auc(fpr, tpr)

# Plotting ROC Curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='darkgray', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

roc_auc

```



0.7788792987423027

Quantitative Analysis

```
import statsmodels.api as sm

# Reconstructing the logistic model using statsmodels for detailed statistics
# Adding a constant to the features for the intercept
X_train_sm = sm.add_constant(X_train)

# Building the logistic model using statsmodels
logit_model = sm.Logit(y_train, X_train_sm)
result = logit_model.fit()

# Displaying the summary of the logistic regression model
model_summary = result.summary2()
model_summary
```

```
Optimization terminated successfully.
Current function value: 0.431349
Iterations 7
```

```
Model:          Logit          Method:          MLE
Dependent Variable: churn      Pseudo R-squared: 0.151
Date:           2023-11-18 19:57 AIC:           6925.5774
No. Observations: 8000        BIC:           7009.4237
Df Model:       11            Log-Likelihood: -3450.8
Df Residuals:   7988          LL-Null:         -4063.5
Converged:      1.0000        LLR p-value:    5.5571e-256
No. Iterations: 7.0000        Scale:          1.0000
```

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
const	-0.8403	2089977.5504	-0.0000	1.0000	-4096281.5677	4096279.8870
x1	-0.0678	0.0302	-2.2454	0.0247	-0.1269	-0.0086
x2	0.7551	0.0300	25.1575	0.0000	0.6963	0.8139
x3	-0.0427	0.0301	-1.4198	0.1557	-0.1018	0.0163
x4	0.1609	0.0357	4.5059	0.0000	0.0909	0.2308
x5	-0.0606	0.0305	-1.9914	0.0464	-0.1203	-0.0010
x6	-0.0103	0.0301	-0.3410	0.7331	-0.0694	0.0488
x7	-0.5341	0.0321	-16.6386	0.0000	-0.5970	-0.4712
x8	0.0158	0.0304	0.5188	0.6039	-0.0438	0.0753
x9	-0.5703	nan	nan	nan	nan	nan
x10	0.2089	nan	nan	nan	nan	nan
x11	-0.4790	nan	nan	nan	nan	nan
x12	-0.1540	nan	nan	nan	nan	nan
x13	-0.6863	nan	nan	nan	nan	nan

```
# Calculating Odds Ratios and their 97.5% Confidence Intervals
odds_ratios = pd.DataFrame()
odds_ratios['Coefficient'] = result.params
odds_ratios['Odds Ratio'] = np.exp(result.params)
odds_ratios['P-value'] = result.pvalues
odds_ratios['[0.025 CI]'] = np.exp(result.conf_int().iloc[:, 0])
odds_ratios['[0.975 CI]'] = np.exp(result.conf_int().iloc[:, 1])

odds_ratios
```